



CONFIDENTIAL

SPARK · ENGINEERING

android developer intern.

take-home assignment. about four hours of work, forty-eight hours to send it back. we are measuring what you can ship in four hours, not what you can polish in four days.

ROLE

android developer intern

PRODUCT

Tiffin · practice product

TIME

~4 hours · 48 hours to submit

SPARK · TAKE-HOME

before you start

Propel ships android features from idea to production in days, not sprints. this assignment is shaped like a normal week here: a spec on monday, something installable by wednesday, and an honest account of what you'd fix next.

this assignment is about **four hours of work** and you have **48 hours** from when you receive it. it is deliberately more than four hours of ideal work — cutting the right scope is part of what we're reading.

the product you're working with

Tiffin is a practice product. we invented it for this assignment — it is not one of ours and there is nothing to look up.

an android app for ordering a daily tiffin from home kitchens nearby. free users browse; paid users get delivery slots and the weekly menu. the money works like the rest of our portfolio: a ₹1 trial purchase, then 24 hours later the real plan activates on its own at ₹249 a month.

ONE THING WORTH KNOWING

if something about the spec isn't clear, ask us. asking a good question is a point in your favour, not a penalty. guessing silently is the thing we're watching for.

ON AI

use Claude, Cursor, Copilot, whatever you work with. that is how this team writes code, and doing this by hand to look impressive would be the wrong signal. what we're reading is whether you can direct an agent and catch it when it's wrong — task 2 and task 4 are entirely about that.

task 1 — ship it · 150 MIN

build a single-module Compose app, kotlin, min sdk 24. no backend to call — bundle the data as a local json asset you write yourself, twelve kitchens is plenty.

three screens:

- **list** — kitchens near you: name, cuisine, price per tiffin, veg/non-veg, rating. real loading, empty and error states, not a spinner that never resolves.
- **detail** — one kitchen, its weekly menu, and a "subscribe" button.
- **paywall** — opens on the third launch and on any subscribe tap. it must state, in words, that ₹1 is charged now and ₹249 on a named date 24 hours later. no payment sdk — fake the purchase, but persist it, and don't show the paywall again once "paid".

it must survive rotation and process death, and the state must still be right after force-stopping and reopening. ship an `analytics.kt` with a single interface and a logcat implementation, called at the four moments you think matter — we'll read which four you picked.

SCOPE IS THE TEST

you will not finish everything in 150 minutes. decide what to cut, cut it deliberately, and put a short "what i cut and why" at the top of your readme. an honest cut beats a half-finished screen every time.

task 2 — fix what the agent wrote · 30 MIN

this compiles and runs. an agent produced it for the list screen. it has four separate problems — one of them will get us a one-star review, one of them will get us a bill.

```
@Composable
fun KitchenList(vm: KitchenViewModel) {
    var items = mutableListOf<Kitchen>()
    LaunchedEffect(items) {
        val res = URL(BASE + "?key=" + "sk_live_9f2c41ab")
            .readText()
        items.addAll(parse(res))
    }
    LazyColumn { items(items) { KitchenRow(it) } }
}
```

name each problem in one line, say what breaks for the user or for us, and paste the corrected version. you do not need to run it.

task 3 — the release note · 20 MIN

half a page, written as if this were going to the play store tomorrow.

- what you'd test by hand before pressing publish, and on which devices.
- the one thing most likely to break in production, and how you'd find out it broke without a user telling you.
- if it did break at 11pm, what's your rollback?
- the next three things you'd build, in order.

task 4 — how you used ai · 15 MIN

- which tools, and roughly what share of the code they wrote.
 - your single best prompt, pasted exactly as you sent it.
 - **one thing the agent got wrong in your own build** — not in task 2, in your actual work — how you caught it, and what you did. an answer saying it was right about everything reads as either no ai or no checking.
 - one thing you decided to write by hand because directing an agent was slower.
-

how to send it back

reply to the same email thread you got this on with:

- a github link — public, or private with propel-hiring invited. commit as you go; we will look at the history.
- a debug apk we can install, attached or linked.
- a readme carrying your cut list, task 2 and task 3.
- a screen recording under 90 seconds, phone or emulator, showing the three screens and the paywall not coming back after payment.

if it doesn't build on our machine we'll tell you and give you one chance to fix it, so a boring readme with exact steps is worth your last five minutes.

CONFIDENTIAL

this document is Propel internal material, shared with you alone for this assessment. please don't forward it, post it, or use it in a portfolio.

the code you write stays yours. we won't ship it or use it in any product.